

Competitive Intelligence Platform: Initial Concept and Deployment Plan

1. Executive Summary

The proposed product is an AI-native competitive and market intelligence platform that helps users build a living intelligence system around the companies, markets, technologies, suppliers, customers, and risks they care about.

The goal is not to build a simple news tracker, scraper, or chatbot wrapper. The goal is to build an evidence-backed intelligence platform that continuously collects external and internal signals, structures them into intelligence objects, connects them through a knowledge graph, and produces role-specific analysis for executives, strategy teams, product teams, sales teams, supply-chain teams, and analysts.

The platform should start as a generic infrastructure with a strong first demonstration vertical. Semiconductors, especially AI semiconductor supply-chain intelligence, is a strong candidate for the first proof vertical because the sector is active, complex, signal-rich, and strategically important.

The long-term product should allow users to create their own intelligence workspace by entering a URL, ticker symbol, company list, sector, or theme. The system should infer the relevant market context, value-chain role, likely competitors, source priorities, signal types, event taxonomy, and recommended monitoring plan.

The initial product philosophy:

Generic engine.

- AI-native industry schema.
- User-defined intelligence universe.
- Evidence-backed analysis.
- Semiconductors as first proof vertical.
- OpenAI-first for POC/MVP.
- Provider-flexible architecture for later.

2. Core Product Concept

The platform should be positioned as:

A configurable intelligence platform where users define the companies, markets, technologies, and sources they care about, and the system turns those signals into an evidence-backed knowledge graph, living dossiers, alerts, and decision-ready analysis.

A simpler positioning line:

Not a news tracker. Not a chatbot. A living intelligence graph for the market you care about.

The system should let a user create a workspace such as:

- AI Semiconductor Supply Chain
- European Medtech Competitors
- Cybersecurity Startup Watchlist
- Industrial Automation Suppliers
- Energy Storage Market Map
- Top 50 Private Robotics Companies

Each workspace should contain:

- Watchlist
- Source policy
- Entity graph
- Signal taxonomy
- Event history
- Evidence library
- Living dossiers
- Analysis outputs
- Alert rules

User feedback memory

The core workflow:

User enters URL, ticker, company list, sector, or theme

- system infers market context
- system suggests competitors, suppliers, customers, technologies, risks, and sources
- user approves or edits the workspace
- system monitors sources
- system extracts evidence and events
- system builds/updates knowledge graph and dossiers
- system generates alerts, memos, battlecards, and executive briefs

3. Why This Should Not Be a Simple CI Wrapper

Many competitive intelligence tools already monitor websites, news, filings, battlecards, and sales enablement content. A new product needs to go deeper.

The weak version would be:

Scrape competitor pages

- summarize with LLM
- send alerts
- let user chat with documents

That is not enough.

The strong version is:

Source registry

- evidence extraction
- entity resolution

- event detection
- source weighting
- confidence scoring
- knowledge graph
- living company/sector dossiers
- RAG over evidence
- strategic analysis
- role-specific outputs

The key difference is that the product should not just tell users what changed. It should explain:

- What changed?
- Why does it matter?
- How confident are we?
- What evidence supports it?
- What related entities are affected?
- What should the user watch next?
- What action should each team consider?

4. Market Reference: What Existing Platforms Teach Us

The platform should learn from the best current products without copying their limits.

AlphaSense-style lesson

AlphaSense demonstrates the importance of deep, trusted content and AI-powered research workflows. The lesson is that source depth, premium research, filings, transcripts, and structured research workflows matter.

What to copy:

- Deep research experience
- Trusted source indexing
- Citation-backed answers
- Company and sector search
- Report generation
- Internal + external content fusion

What to improve:

- Move beyond document search into event intelligence, graph reasoning, and action-oriented workflows.

Klue-style lesson

Klue demonstrates the importance of sales enablement, win/loss intelligence, battlecards, and competitive workflows inside the revenue organization.

What to copy:

- Battlecards
- Deal-specific competitive guidance
- Win/loss feedback
- Sales workflow integration
- Actionable recommendations

What to improve:

- Expand beyond sales into strategy, product, finance, supply chain, and executive decision-making.

Crayon-style lesson

Crayon demonstrates the value of automated competitor monitoring and real-time enablement.

What to copy:

- Competitor monitoring
- Alerts
- Battlecards
- Salesforce/Slack/Teams-style delivery
- Usage metrics

What to improve:

- Reduce noise by turning raw changes into source-weighted, confidence-scored intelligence objects.

Contify-style lesson

Contify demonstrates the value of configurable market and competitive intelligence across markets, competitors, key accounts, partners, and strategic entities.

What to copy:

- Market intelligence orientation
- Custom sources
- Role-specific insights
- Team-level delivery
- Broad strategic entity monitoring

What to improve:

- Make the intelligence graph and event taxonomy central to the product, not just the content feed.

Similarweb-style lesson

Similarweb shows that quantitative digital signals can be powerful in competitive analysis.

What to copy:

- Digital benchmarking
- Traffic/channel insights

- Competitive discovery
- Market share signals
- Digital trend movement

What to improve:

- Connect digital signals to filings, hiring, product changes, patents, regulatory events, customer wins, and strategic implications.

5. AI-Native Industry Context Schema

The platform should not rely on GICS or NAICS as its primary product schema. Those systems can be useful reference points, but they were not designed for AI-native competitive intelligence.

The platform should create its own high-level schema.

Working name:

- AI-Native Industry Intelligence Schema

or:

- Sector Genome Layer

The schema should classify companies and markets across multiple dimensions, not assign each company to one rigid industry bucket.

Each entity or workspace should be profiled across:

- Sector family
- Market arena
- Value-chain role
- Business model
- Product/technology domain
- Buyer segment
- Supplier dependencies
- Competitive basis
- Regulatory exposure
- Source priority model
- Event taxonomy
- Metrics that matter
- Strategic themes

Example profile:

Entity: Nvidia

Sector families:

- Semiconductors
- AI infrastructure

Value-chain roles:

- Fabless chip designer
- Platform/ecosystem provider
- Supplier to hyperscalers

Technology domains:

- GPUs
- AI accelerators
- Networking
- AI software ecosystem

Supplier dependencies:

- Foundry manufacturing
- HBM suppliers
- Advanced packaging
- Substrates

Regulatory exposure:

- Export controls
- China restrictions
- Supply-chain concentration

Event types to monitor:

- Foundry capacity
- HBM supply
- Export-control changes
- Hyperscaler capex
- Competing accelerator launches
- Pricing/margin pressure
- Software ecosystem shifts

This is far more useful than simply assigning a company to an industry code.

6. User Onboarding

The product should make onboarding feel simple.

The user should be able to start with one of the following:

- Ticker symbol
- Company URL
- Company name
- List of companies
- Sector name
- Market theme
- Uploaded spreadsheet

Example:

User enters:

NVDA, AMD, TSMC, Samsung, SK hynix, Micron, ASML

System proposes:

Workspace: AI Semiconductor Supply Chain

Suggested themes:

- AI accelerators
- HBM
- Advanced packaging
- Foundry diversification
- Export controls
- Hyperscaler capex
- Equipment constraints
- Substrates
- Photonics

Suggested sources:

- SEC filings
- Earnings calls
- Investor presentations
- Company websites
- Reuters/AP/financial news
- SEMI and trade sources
- Patent databases
- Government export-control notices
- Job postings
- Supplier announcements

The user should not have to design the whole intelligence system from scratch. The system should propose the initial market map, and the user should approve, reject, or refine.

7. Core Data Model

The central product object should not be a document. It should be an intelligence object.

An intelligence object should include:

- Event or signal
- Entities involved
- Event type
- Timestamp
- Source
- Evidence snippet

- Source weight
- Confidence score
- Related graph links
- Business implication
- Contradictory evidence
- Recommended watch items

Recommended actions

Example:

Event:

Company X appears to be expanding advanced packaging capacity.

Entities:

- Company X
- Advanced packaging
- AI accelerators
- HBM
- Hyperscaler customers

Evidence:

- Job postings
- Supplier announcement
- Capex language in filing
- Trade publication mention

Confidence:

Medium-high

Implication:

Company X may be positioning for AI infrastructure supply-chain demand.

Recommended watch:

- Customer announcements
- Additional hiring
- Equipment orders
- Capex commentary
- Supplier partnerships

This object-based approach allows the system to build a living intelligence graph instead of a pile of summaries.

8. Knowledge Graph and RAG Strategy

The platform should use both RAG and a knowledge graph.

RAG is useful for finding and summarizing supporting evidence. The knowledge graph is useful for understanding relationships, causality, dependencies, and market structure.

The correct architecture:

Raw source archive

- + normalized evidence store
- + vector index / RAG
- + entity and event database
- + knowledge graph
- + living dossiers
- + analyst outputs

The knowledge graph should connect:

- Company
- Product
- Technology
- Customer
- Supplier
- Partner
- Executive
- Investor
- Regulator
- Geography
- Facility
- Patent
- Filing
- Event
- Claim
- Risk
- Opportunity

Example graph relationships:

- Company → supplies → Customer
- Company → depends_on → Supplier
- Company → manufactures_at → Facility
- Company → competes_with → Company
- Product → uses → Technology
- Technology → constrained_by → Material

- Event → supported_by → Evidence
- Risk → affects → Company

The graph should not be overbuilt in the first version. The MVP can start with a simple property graph model in Postgres or a graph database if the team is comfortable. The key is to structure entities and relationships from day one so the product can grow into a richer graph later.

9. Living Dossier Layer

The platform should maintain living dossiers for companies, sectors, technologies, and themes.

A living dossier is an AI-maintained, evidence-backed profile that updates as new information arrives.

Examples:

- /company/nvidia
- /company/tsmc
- /sector/ai-semiconductor-supply-chain
- /theme/hbm-bottleneck
- /theme/export-controls
- /theme/advanced-packaging

Each dossier should include:

- Current strategy
- Recent moves
- Products and technology areas
- Customers and suppliers
- Competitive strengths
- Competitive weaknesses
- Risk exposure
- Financial/operating signals
- Recent evidence
- Open questions
- Watch items

The living dossier should not replace the evidence store, RAG, or knowledge graph. It should be the readable layer that makes the system easy for humans to understand.

10. Source Strategy

Sources should be ranked by importance. The platform should not treat all sources equally.

Generic source tiers:

Tier 0: Authoritative primary sources

Examples: SEC filings, government/regulatory data, court filings, patent offices, official procurement records

Tier 1: Trusted licensed news and research

Examples: Reuters, AP, Bloomberg, FT, WSJ, industry research, premium databases

Tier 2: Company-controlled sources

Examples: websites, investor presentations, press releases, product pages, pricing pages, job postings

Tier 3: Sector-specific sources

Examples: trade journals, conference programs, standards bodies, technical publications

Tier 4: Market/digital signals

Examples: web traffic, SEO, app data, review sites, ad signals, social mentions

Tier 5: Internal field intelligence

Examples: CRM notes, sales calls, win/loss notes, customer support, partner feedback

Every extracted intelligence object should carry source metadata:

- Source name
- Source type
- Source tier
- Published date
- Collected date
- Reliability
- Bias risk
- License status
- Evidence URL or internal reference

This creates trust and auditability.

11. First Proof Vertical: Semiconductors

The infrastructure should be generic, but the first proof vertical should be narrow and high-value.

Recommended first vertical:

AI Semiconductor Supply Chain Intelligence

Scope:

- AI accelerators
- Foundries
- HBM
- Advanced packaging
- Substrates
- Equipment vendors
- Materials
- Photonics

- Export controls
- Hyperscaler demand
- Fab capacity

Initial company/entity set:

- Nvidia
- AMD
- Broadcom
- Marvell
- TSMC
- Samsung Foundry
- Intel Foundry
- SK hynix
- Micron
- Samsung Memory
- ASML
- Applied Materials
- Lam Research
- KLA
- Tokyo Electron
- ASE
- Amkor
- Ibiden
- Unimicron
- Coherent
- Lumentum
- Google
- Amazon
- Microsoft
- Meta

The strategic question:

Who is gaining or losing position in the AI semiconductor supply chain, and why?

This is better than trying to cover all semiconductors at once.

The semiconductor demo should show:

- Sector map
- Entity graph
- Company dossiers
- Event feed
- Supply-chain risk map
- Weekly intelligence brief
- Source-ranked alerting
- Competitor strategy memo

12. Product Modules

Module 1: Workspace Builder

Allows users to create a market intelligence workspace from URL, ticker, company list, sector, or theme.

Core features:

- Entity input
- Ticker/URL enrichment
- Industry-context inference
- Suggested competitors
- Suggested suppliers/customers
- Suggested technologies/themes
- Suggested source policy
- User approval workflow

Module 2: Source Registry

Defines which sources are monitored and how they are weighted.

Core features:

- Source type
- Source tier
- Refresh frequency
- License/compliance status
- Parser/connector configuration
- Reliability score
- Sector relevance

Module 3: Ingestion and Normalization

Collects and normalizes raw documents and source data.

Core features:

- Web collection
- SEC/filing collection
- News/API collection
- Company website monitoring
- Patent source collection
- Job posting collection
- Document deduplication
- Metadata extraction
- Raw archive

Module 4: Evidence Extraction

Turns documents into evidence records.

Core features:

- Entity extraction
- Claim extraction
- Event extraction

- Evidence snippets
- Source citation
- Timestamp extraction
- Confidence scoring
- Contradiction detection

Module 5: Intelligence Object Engine

Converts evidence into structured intelligence objects.

Core features:

- Event classification
- Entity linking
- Source weighting
- Confidence scoring
- Impact hypothesis
- Related entities
- Recommended watch items
- Analyst review status

Module 6: Knowledge Graph

Connects entities, events, claims, and evidence.

Core features:

- Company relationships
- Customer/supplier relationships
- Product/technology relationships
- Market relationships
- Risk exposure relationships
- Event-to-evidence relationships

Module 7: RAG and Search

Lets users search and ask questions grounded in evidence.

Core features:

- Hybrid search
- Vector retrieval
- Keyword retrieval
- Source filters
- Evidence-backed answers
- Citation display
- Workspace-specific retrieval

Module 8: Living Dossiers

Maintains current profiles for companies, sectors, technologies, and themes.

Core features:

- Company profiles
- Sector profiles
- Theme profiles
- Recent changes
- Open questions
- Watch items
- Evidence trail
- Dossier refresh workflow

Module 9: Analysis and Output Engine

Produces role-specific intelligence outputs.

Core features:

- Weekly brief
- Executive memo
- Board memo
- Sales battlecard
- Product threat memo
- Supplier risk report
- Market-entry analysis
- Competitor strategy memo

Module 10: Feedback and Learning

Allows users to correct the system and improve workspace relevance.

Core features:

- Approve/reject event
- Correct entity link
- Adjust source priority
- Add/remove themes
- Tune alert sensitivity
- Save preferred output formats

13. Deployment Architecture

The POC/MVP should be cloud-based, API-first, and simple. Avoid premature self-hosting of open models.

Recommended POC/MVP stack:

Frontend:

- React / Next.js

Backend:

- Python FastAPI or Node.js/TypeScript
- Database:
- Postgres
- Vector search:
- pgvector to start, or managed vector DB if preferred

- Object storage:
- S3-compatible storage for raw documents and normalized artifacts

Queue/workers:

- Celery/RQ, Sidekiq, BullMQ, or managed queue

Scheduler:

- Cron/Temporal/Airflow-lite depending on team preference

LLM:

- OpenAI API
- Embeddings:
- OpenAI embeddings initially

Knowledge graph:

- Start with Postgres graph-style tables or Neo4j if graph expertise exists

Auth:

- Basic workspace/user auth for MVP

Deployment:

- AWS, GCP, or Azure; AWS is a practical default
- A simple MVP architecture:

User workspace

- → source registry
- → ingestion workers
- → raw document store
- → parser/normalizer
- → evidence extraction
- → event/intelligence object store
- → vector index
- → knowledge graph
- → dossier generator
- → analysis/output engine
- → UI + alerts

14. Model Strategy

Use OpenAI through POC and MVP.

Do not self-host Gemma or another open-weight model initially unless there is a hard privacy, offline, or volume requirement.

Recommended model usage:

Cheaper/smaller model:

- page classification
- source tagging
- simple entity extraction
- routing documents to event types
- deduplication
- short summaries

Stronger model:

- strategic analysis
- contradiction detection
- event implication
- competitor strategy inference
- board memo generation
- sector context generation
- user-facing synthesis

Embeddings:

- use OpenAI embeddings initially

Batch jobs:

- use asynchronous/batch processing for overnight refreshes

Important design principle:

- Do not send giant raw documents to the model repeatedly.
- Convert sources into compact evidence objects first.
- Then reason over structured evidence.

The architecture should include a model router from the beginning:

Task type

- + latency need
- + cost budget
- + quality requirement
- + privacy requirement
- + context size
- → model choice

This keeps the platform flexible if the team later moves some commodity tasks to open models.

15. POC Scope

The POC should prove that the system can produce differentiated intelligence from a small but meaningful corpus.

Recommended POC:

Vertical:

AI Semiconductor Supply Chain

Entities:

25–50 companies/entities

Sources:

- SEC filings
- Company websites
- Investor presentations
- Press releases
- Selected news sources
- Selected sector sources
- Patent/job-posting data if feasible

Outputs:

- Company dossier
- Sector map
- Event feed
- Weekly intelligence brief
- One competitor strategy memo
- One supplier risk memo

POC success criteria:

- Can the system classify companies correctly from ticker/URL?
- Can it suggest useful related entities?
- Can it extract meaningful events?
- Can it link evidence to events?
- Can it generate a useful company dossier?
- Can it generate a weekly brief that a human would actually read?
- Can the user trace every claim back to evidence?

16. MVP Scope

The MVP should support real users creating and managing workspaces.

Recommended MVP features:

- User accounts and workspaces
- Ticker/URL/company-list onboarding
- AI-native industry context inference
- Source registry
- Basic source monitoring
- Evidence extraction
- Event/intelligence object creation
- Entity graph

- Vector search
- Living dossiers
- Weekly brief
- Alert feed
- User feedback controls
- Exportable memo/battlecard

Avoid in MVP:

- Too many verticals
- Too many source types
- Complex graph UI
- Self-hosted model infrastructure
- Heavy enterprise integrations
- Full CRM/Gong/Slack integrations

17. Roadmap

Phase 0: Design

- Define product thesis
- Define AI-native industry schema
- Define source tier model
- Define intelligence object schema
- Define semiconductor proof vertical
- Create mock UI flows

Phase 1: POC

- Build ingestion for limited source set
- Build entity extraction
- Build event extraction
- Build evidence store
- Build basic vector search
- Build simple graph structure
- Generate company dossiers and weekly brief

Phase 2: MVP

- Workspace creation
- Ticker/URL onboarding
- User approval workflow
- Source registry UI
- Event feed
- Dossier UI
- Search/Q&A
- Weekly brief
- Feedback loop

Phase 3: Pilot

- Measure usefulness of briefs and alerts

- Tune event taxonomy and source weighting
- Add role-specific outputs

Phase 4: Expansion

- Add new vertical packs
- Add private/internal sources
- Add Slack/Teams/email delivery
- Add CRM/sales enablement workflows
- Add advanced graph analytics
- Add model cost optimization
- Consider open models for commodity tasks

18. Key Risks

Risk 1: Product becomes a noisy alert feed

Mitigation:

Every alert must have source weighting, confidence, business implication, and recommended watch items.

Risk 2: Product becomes a generic chatbot

Mitigation:

Center the product around intelligence objects, graph relationships, dossiers, and role-specific outputs.

Risk 3: Classification schema becomes too academic

Mitigation:

Keep the schema practical. It must help onboarding, source selection, event detection, and output quality.

Risk 4: Source licensing issues

Mitigation:

Use official APIs and licensed sources where required. Avoid scraping restricted sources. Track license status in the source registry.

Risk 5: Open-ended vertical scope

Mitigation:

Start with one sharp proof vertical: AI semiconductor supply-chain intelligence.

Risk 6: Model costs grow unexpectedly

Mitigation:

Use compact evidence objects, caching, batch jobs, smaller models for commodity tasks, and a model router.

Risk 7: Trust issues

Mitigation:

Every claim should link to evidence. Show confidence and contradictory evidence when available.

20. Initial Build Priorities

The first build should focus on five things:

1. AI-native industry context schema
2. Workspace onboarding from ticker/URL/company list
3. Source registry and source weighting
4. Evidence → event → intelligence object pipeline
5. Living dossier and weekly brief generation

If these five work, the product has a strong foundation.

The first demo should show:

- User enters a semiconductor company list.
- System builds a workspace.
- System suggests related themes and entities.
- System ingests selected sources.
- System extracts events.
- System creates company dossiers.
- System generates a weekly intelligence brief.
- User can click every claim and see supporting evidence.

21. Recommended Initial Positioning

Main positioning:

A configurable intelligence platform that builds a living, evidence-backed market graph around the companies, technologies, suppliers, customers, and risks you care about.

Sharper version:

Most CI tools tell you what changed. This tells you what the change means, how confident you should be, and what action each team should consider.

Technical positioning:

An AI-native competitive intelligence system using source-ranked evidence, event extraction, knowledge graphs, RAG, and living dossiers.

First vertical positioning:

AI semiconductor supply-chain intelligence for strategy, sales, and supply-chain teams.

22. Bottom Line

The strongest product direction is:

Generic infrastructure

- + AI-native industry schema
- + user-defined intelligence workspaces
- + source-ranked evidence

- + structured intelligence objects
- + knowledge graph
- + RAG
- + living dossiers
- + role-specific analysis
- + OpenAI-first POC/MVP
- + semiconductor proof vertical

This gives the product both flexibility and depth.

The mistake would be building a generic scraper and chatbot. The opportunity is to build a living intelligence system that understands markets, tracks evidence, detects meaningful events, and turns signals into decisions.